

Hoang

MATH158 Final Project

December 7th, 2023



Predicting Used Value

What's the maximum listing price my used car could be sold?

- 3,000,000 observations from (kaggle.com/us-used-cars-dataset) – 10GB.
- 66 features, half of them have no descriptions in metadata.
- Crawled on cargurus.com in September 2020 -> submissions are **independent**.
- Diverse data types: text, number, Nan
- Metric: custom Mean Absolute Error, where:
 - under-prediction will * 1.2
 - over-prediction will * 0.8

=> **Motivation Questions:**

1. What features should be used for prediction?
2. How well will my prediction be **without knowing any feature descriptions?**

Explore Data: getting started

- What's my response column name? "price".

=> X: DataFrame without column 'price'.

y: Column 'price' and nothing else.

- Do I want an intercept β_0 ?

<= Are there rows where all features' values being zero(s) ?

```
> sum(apply(X, 1, function(row) all(row == 0)))
```

```
[1] 0
```

=> YES

=> X: DataFrame without column 'price' and has a column full of 1's.

- Do I want duplicating rows / columns ?

<= Will duplicating rows / columns affect regression coefficients? No.

<= Will duplicating rows / columns invalidate regression assumptions? Yes (multi-collinearity).

=> NO, remove them.


- How do I split the data?

=> Just split as usual (80% X_train, y_train; 20% X_test, y_test)

Explore Data: welcoming “the null”

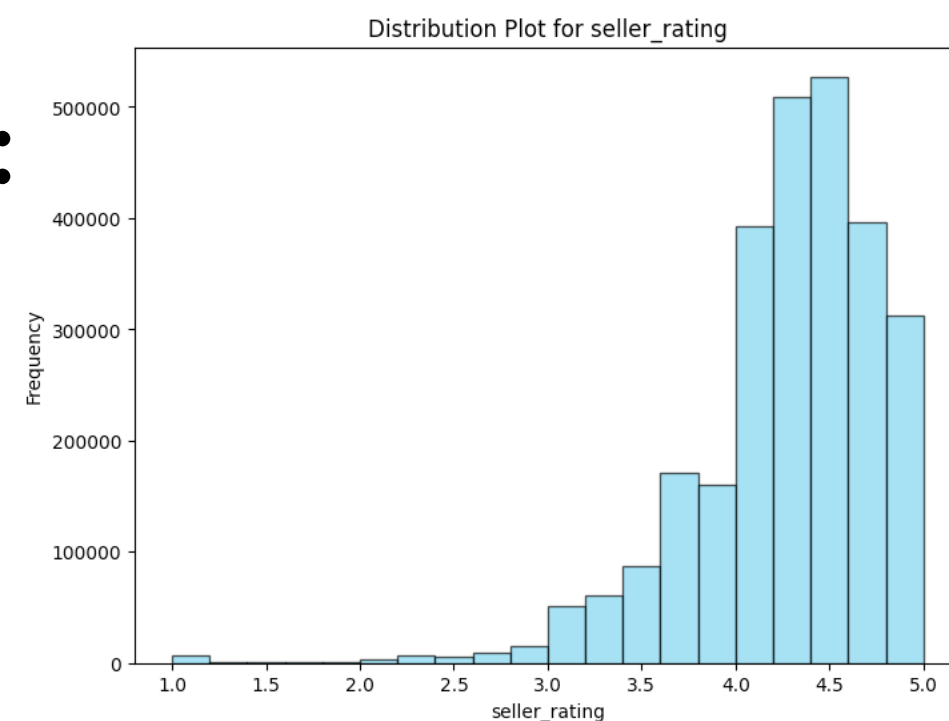
- Remove “massive null” columns having 20%+ (600,000+) missing observations.
- Impute “small null” columns: based on the distribution of a column

- If this:



=> Impute mean

- If that:




=> Impute median

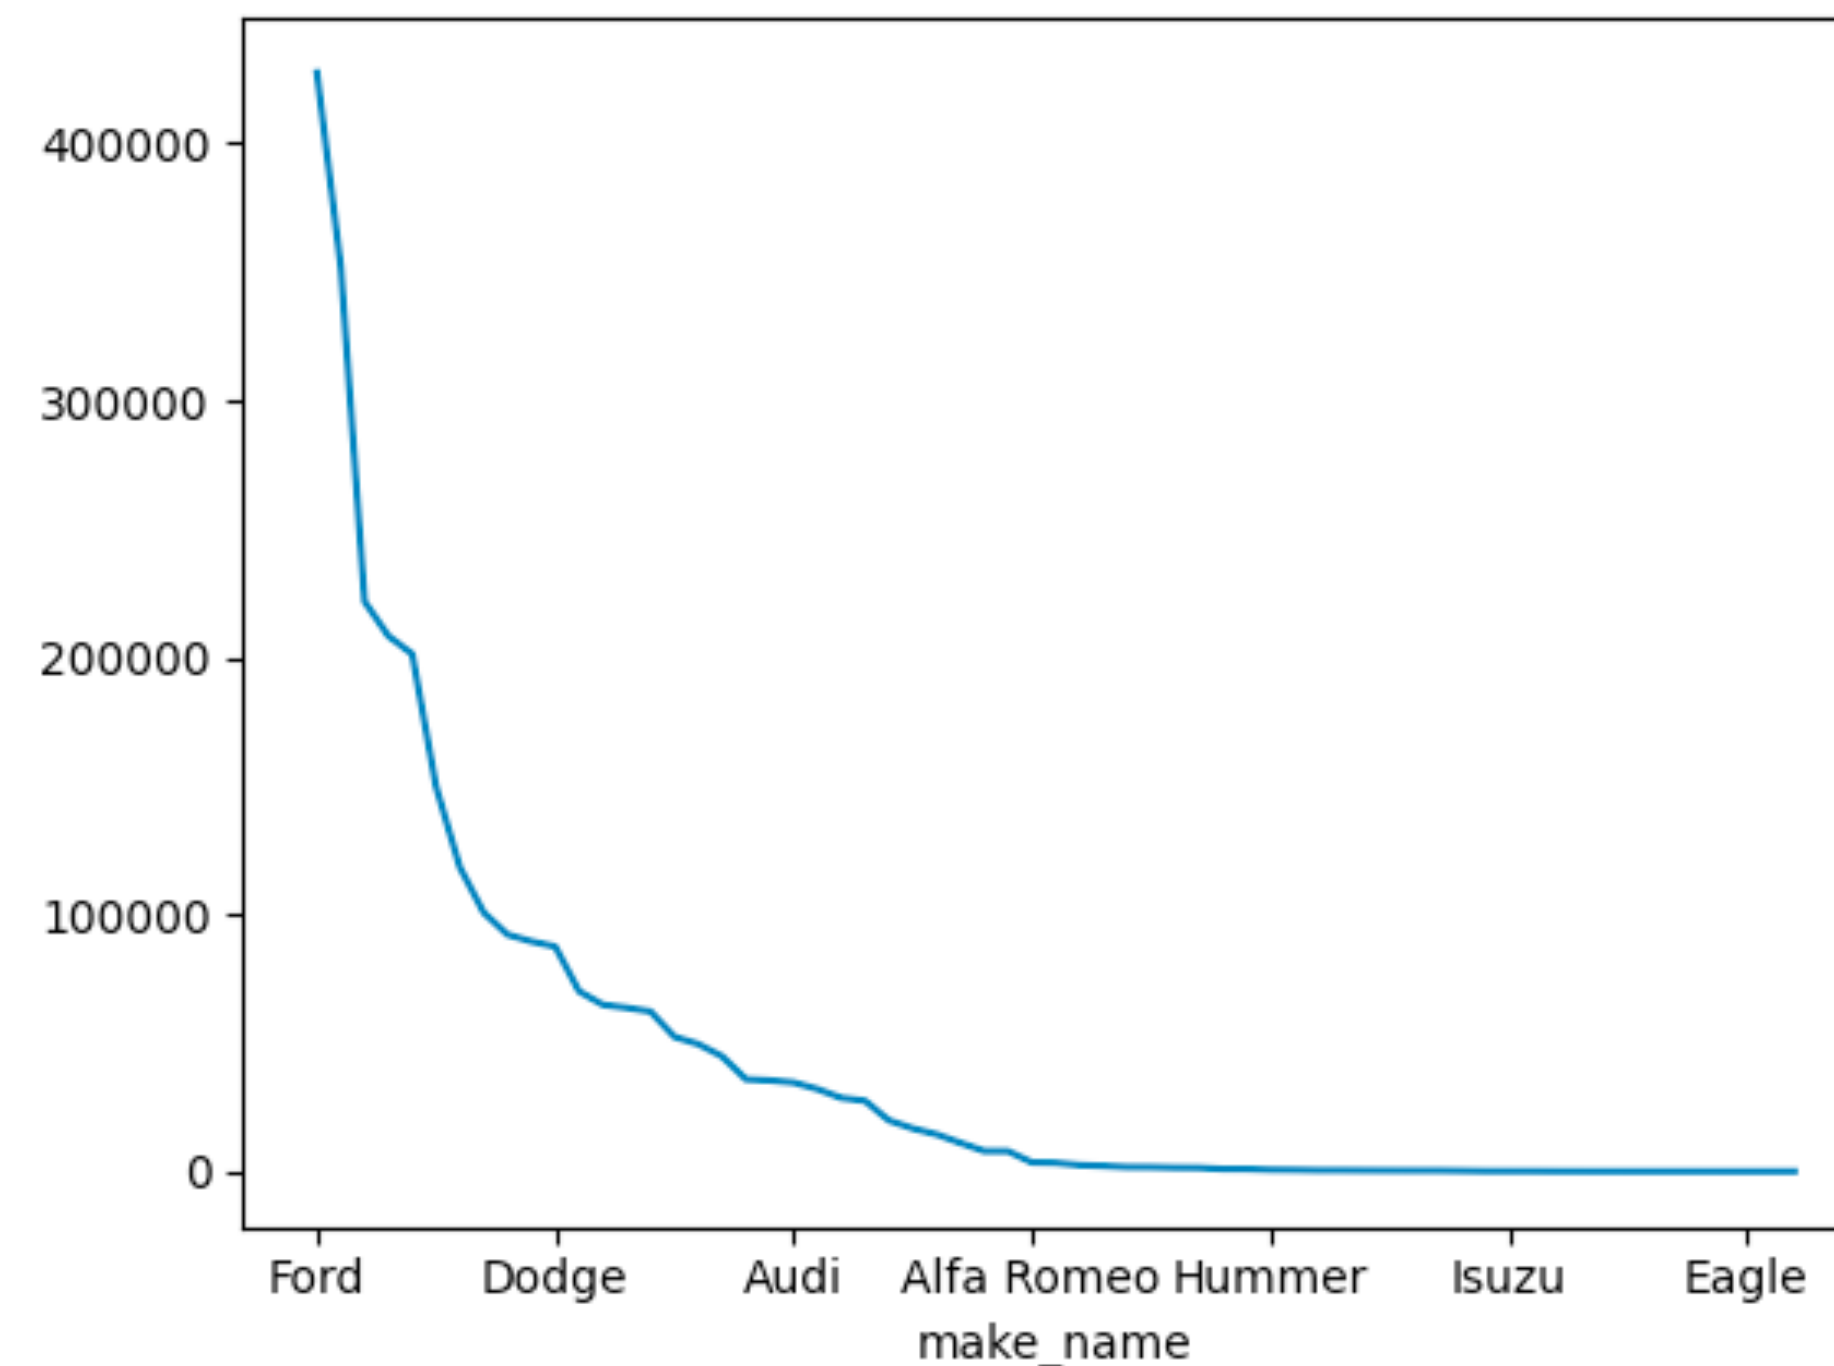
- Else:

=> Impute mode

Explore Data: translating “the text”



What if we're overcrowded !?




- What do I want from encoding? 1 to 2 same meaning as 2 to 3 \leq IQR Frequency Encoding.

The Full Linear Model: preparing the ingredient

- Do I want to transform y?

=> YES



- Do I want to transform X? <= When do I need to transform X?

- Non-Linearity magic Linearity
- Something needs it (PCA, PLS, Ridge, ...)

=> JUST IN CASE

=>

```
get_standardized_df <- function(df) {  
  # if (a column is non-encoded) {  
  #   standardize(column)  
  # }  
}
```

 => std_X_train, std_y_train, std_...

The Full Linear Model: let it cook



- What's the R-Squared(s)?

Residual standard error: 10610 on 2201001 degrees of freedom
Multiple R-squared: 0.674, Adjusted R-squared: 0.674
F-statistic: 7.583e+04 on 60 and 2201001 DF, p-value: < 2.2e-16

- How long did it cook?

user 37.560 system.time({ *your_full_linear_model_here* })

- Does it violate normality and constant variance assumptions?



=> YES 🥲

- Does multi-collinearity exist?

=> YES

```
> vif(full_linear_model)
Error in vif.default(full_linear_model) :
  there are aliased coefficients in the model
```

🚩	Linearity
🚩	Independence
🚩	Normality
🚩	Constant Variance

The First Red Flag: Multi-collinearity

I want: VIF(a linear model) < 5, ideally ~1 to 2

Principle Component Analysis

Partial Least Squares

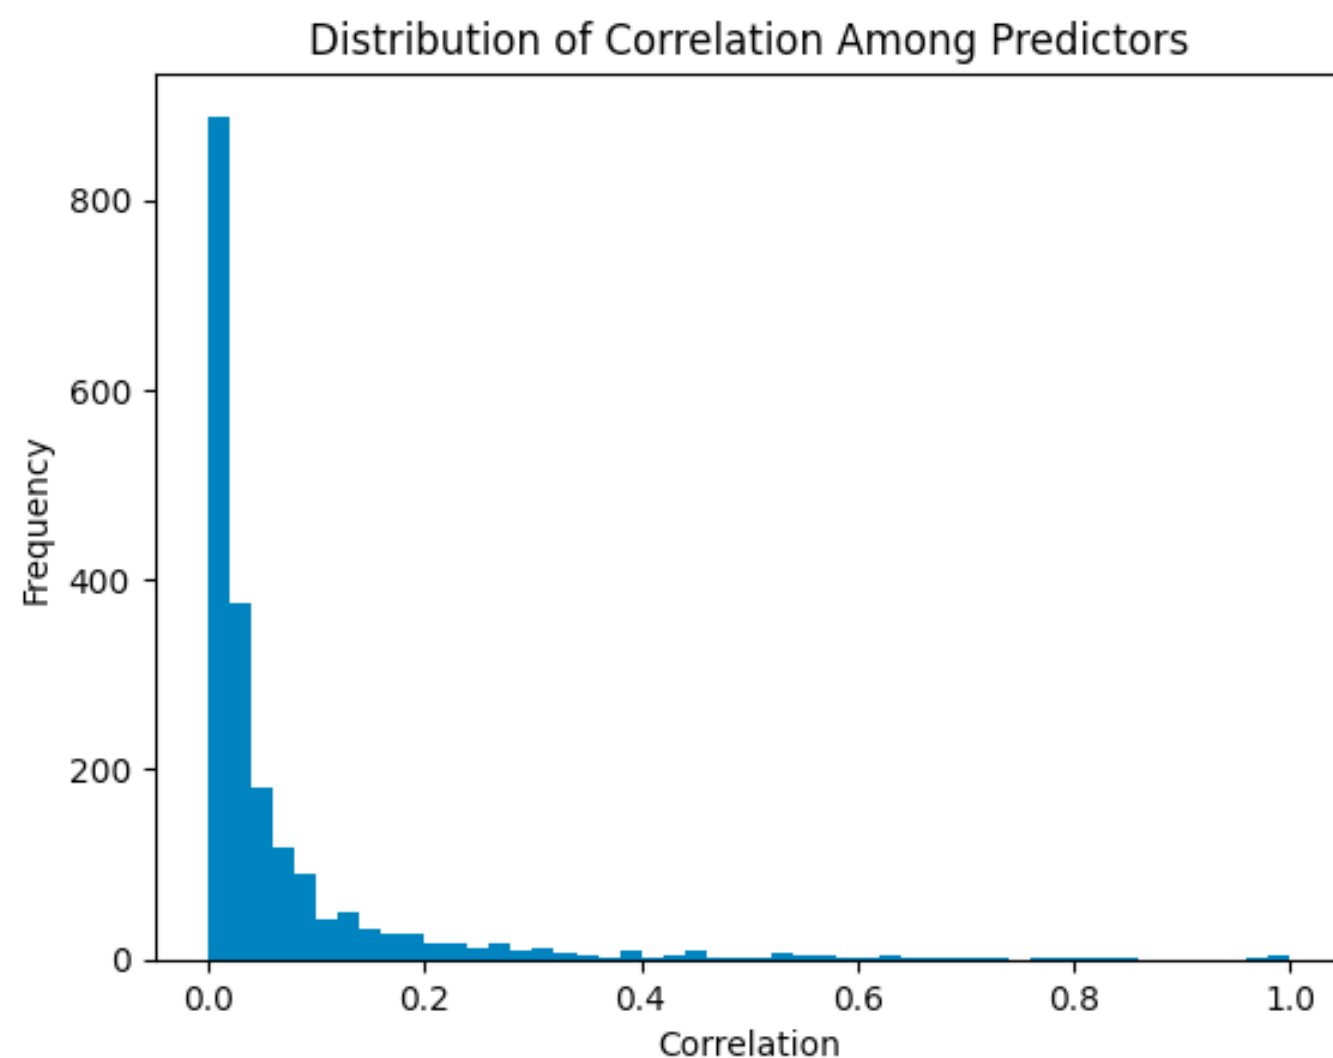
Ridge Regression

If majority of correlations among predictors are small, PCA is not effective

If majority of correlations with response are small, PLS is not effective

1. $d(\text{ridge function}) = -2y^T X + 2X^T X \hat{\beta} + 2\lambda \hat{\beta}$
2. Solution when $d(\text{ridge function}) = 0$:
$$\hat{\beta} = y^T x / (x^T x + \lambda)$$


always > 0 whatever the penalty rate



```
##{r}  
correlation_with_y <- cor(df)  
correlation_without_itself <- correlation_with_y[, 'price'][-which(names(correlation_with_y[, 'price']) == 'price')]  
mean(correlation_without_itself)  
##  
[1] 0.03499676
```

The LASSO
Solution when $d(\text{lasso function}) = 0$:
1. $\hat{\beta} = (y^T x - \lambda) / (x^T x)$ when $\lambda \leq y^T x$
2. $\hat{\beta} = (y^T x + \lambda) / (x^T x)$ when $\lambda \geq y^T x$

The LASSO
then OLS



Whatever_Error(LASSO)
<
Whatever_Error(OLS)

The First Red Flag: Multi-collinearity

I want: $VIF(\text{a linear model}) < 5$, ideally ~ 1 to 2 , find how much should I penalize for the LASSO

- Idea 1:
 - generate a list of penalties (lambdas), fit models, find the lambda with lowest $\max(VIF)$.

```
```{r}
find_optimal_lambda <- function(lambda_options, X_train, y_train) {
 # Standardize X_train, y_train
 for (lambda_option in lambda_options) {
 lasso_model = glmnet(std_X_train, std_y_train, alpha = 1, lambda = lambda_option)

 # Filter non-zero coefficient columns
 non_zero_rows = rownames(coef_values)[coef_values[, 's0'] != 0]
 lasso_stayed_columns <- non_zero_rows[2:length(non_zero_rows)] # Avoid (Intercept)
 # \Filter non-zero coefficient columns

 new_linear_model <- lm(y_train ~ ., data = X_train)
 cur_vif = vif(new_linear_model)

 # Update min_lambda if there is new min_vif
 }
}
lambda_options = seq(from = 0, to = 2, length.out = 100)
```
```



R Session Aborted

R encountered a fatal error.
The session was terminated.


[Start New Session](#)

- There is a relationship between lambda and VIF => Can we compute VIF from lambda?

The First Red Flag: Multi-collinearity

I want: $VIF(\text{a linear model}) < 5$, ideally ~ 1 to 2 , find how much should I penalize for the LASSO

- Idea 2:
 - $VIF = 1 / (1 - R^2)$
 - R^2 is how much variance of actual_y (represented by SST) are from predicted_y (SSReg)
 - Geometrically:



$$\Rightarrow SSReg = \text{cov_mat}(y, X)^T @ \text{cov_mat}(X) @ \text{cov_mat}(y, X)$$

$$\Rightarrow SST = \text{cov_mat}(y)$$

$$\Rightarrow R^2 = SSReg / SST = SSReg @ SST^{-1}$$

$$= \text{cov_mat}(y, X)^T @ \text{cov_mat}(X)^{-1} @ \text{cov_mat}(y, X)$$

$$\Rightarrow VIF = 1 / (1 - R^2) \Rightarrow \text{cov_mat}(X)^{-1} = F(\text{cov_mat}(X))$$

$$\Rightarrow VIF_{\text{regularization}} = F(\text{cov_mat}(X) + \text{lambda} * I)$$

The First Red Flag: Multi-collinearity

I want: $VIF(\text{a linear model}) < 5$, ideally ~ 1 to 2 , find how much should I penalize for the LASSO

- Idea 2: Proof

```
def vif_statsmodels(X, intercept_colname):
    from statsmodels.stats.outliers_influence import variance_inflation_factor
    from statsmodels.tools.tools import add_constant

    X_copy = X.copy()
    if intercept_colname not in X_copy: X_copy = add_constant(X_copy)

    vif = pd.DataFrame()
    vif["variables"] = X_copy.columns
    vif["VIF"] = [(variance_inflation_factor(X_copy.values, i), 4) for i in range(X_copy.shape[1])]
    return vif

def vif_custom(X, lam):
    from numpy.linalg import inv

    n_cols = X.shape[1] - 1
    vif = np.zeros((len(lam), n_cols))

    rxx = X.iloc[:, :n_cols].corr().values
    rxy = X.corr().iloc[:n_cols, n_cols].values
    for i in range(len(lam)):
        tmp1 = inv(rxx + lam[i] * np.eye(n_cols))
        vif[i, :] = np.diag(np.dot(np.dot(tmp1, rxx), tmp1))

    vif_df = pd.DataFrame(vif, columns=X.columns[:-1], index=lam)

    return vif_df.T
```

```
vif_df.T.head()
```

✓ 0.0s

| 0 | |
|---------------------|----------|
| back_legroom | 1.583718 |
| city | 1.085481 |
| daysonmarket | 1.163538 |
| engine_displacement | 6.345453 |
| engine_type | 1.990666 |

```
vif_statsmodels_df.head()
```

✓ 0.0s


| | variables | VIF |
|---|---------------------|------------|
| 0 | const | 530.271953 |
| 1 | back_legroom | 1.583747 |
| 2 | city | 1.085473 |
| 3 | daysonmarket | 1.163727 |
| 4 | engine_displacement | 6.345435 |

The First Red Flag: Multi-collinearity

I want: $VIF(\text{a linear model}) < 5$, ideally ~ 1 to 2 , find how much should I penalize for the LASSO

- Idea 2:
 - generate a list of penalties (lambdas), find the smallest lambda that VIF table has > 0 rows, collect all the rows, fit each row with a linear model, then pick the model having the lowest AIC.

Result for `optimal_lambda = 0.0440044`



Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|----------------------------------|------------|------------|---------|------------|
| (Intercept) | -2.788e+02 | 4.449e-01 | -626.73 | <2e-16 *** |
| back_legroom | 1.225e-02 | 1.484e-04 | 82.54 | <2e-16 *** |
| engine_type | -2.556e-01 | 1.702e-03 | -150.17 | <2e-16 *** |
| fuel_tank_volume | 2.041e-02 | 1.458e-04 | 139.95 | <2e-16 *** |
| horsepower | 8.130e-03 | 8.112e-06 | 1002.18 | <2e-16 *** |
| longitude | -2.943e-03 | 2.983e-05 | -98.65 | <2e-16 *** |
| make_name | -1.307e-01 | 1.809e-03 | -72.24 | <2e-16 *** |
| mileage | -1.209e-05 | 1.816e-08 | -665.75 | <2e-16 *** |
| model_name | -1.038e-01 | 1.690e-03 | -61.45 | <2e-16 *** |
| savings_amount | 9.013e-05 | 4.573e-07 | 197.10 | <2e-16 *** |
| seller_rating | 1.320e-01 | 8.086e-04 | 163.19 | <2e-16 *** |
| width | 1.897e-03 | 7.591e-05 | 25.00 | <2e-16 *** |
| year | 1.462e-01 | 2.214e-04 | 660.61 | <2e-16 *** |
| is_body_type_Sedan | -1.442e-01 | 1.054e-03 | -136.77 | <2e-16 *** |
| `is_body_type_Pickup Truck` | -3.103e-01 | 1.548e-03 | -200.46 | <2e-16 *** |
| is_body_type_Convertible | 4.778e-01 | 4.583e-03 | 104.24 | <2e-16 *** |
| is_fuel_type_Diesel | 5.379e-01 | 4.236e-03 | 126.97 | <2e-16 *** |
| `is_fuel_type_Flex Fuel Vehicle` | -3.865e-01 | 2.003e-03 | -192.96 | <2e-16 *** |
| is_fuel_type_Hybrid | 4.225e-01 | 2.614e-03 | 161.62 | <2e-16 *** |
| is_is_new_False | -2.910e-01 | 1.245e-03 | -233.71 | <2e-16 *** |
| is_franchise_dealer_False | -2.183e-01 | 1.262e-03 | -172.91 | <2e-16 *** |
| is_wheel_system_FWD | -4.654e-01 | 1.046e-03 | -444.75 | <2e-16 *** |
| is_wheel_system_4X2 | -2.686e-01 | 2.190e-03 | -122.64 | <2e-16 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6067 on 2201039 degrees of freedom

Multiple R-squared: 0.861, Adjusted R-squared: 0.861

F-statistic: 6.199e+05 on 22 and 2201039 DF, p-value: < 2.2e-16

The Second Red Flag: Normality

- Should I use Robust Regression?
- \leq How much percentage of outliers are there and are most of them influential?

```
abnormal_points = get_abnormal_points(lassoAIC_linear_model)
influentials <- abnormal_points$influential_points

outliers <- abnormal_points$outliers
nonInfluential_outliers = setdiff(influentials, outliers)
influential_outliers = intersect(influentials, outliers)

leverages <- abnormal_points$leverage_points
nonInfluential_leverages = setdiff(influentials, leverages)
influential_leverages = intersect(influentials, leverages)

all_abnormal_points <- union(outliers, leverages)

cat(length(all_abnormal_points) * 100 / dim(lassoAIC_X_train)[1], "\n")
cat(length(influentials) * 100 / dim(lassoAIC_X_train)[1], "\n")
cat(length(nonInfluential_outliers) * 100 / length(outliers), "\n")
cat(length(nonInfluential_leverages) * 100 / length(leverages), "\n")
````
```

12.62141  
6.542206  
74.95397  
27.23772



$\Rightarrow$

```
````{r}
removed_lassoAIC_X_train <- lassoAIC_X_train[!(rownames(lassoAIC_X_train) %in% influentials), ]
removed_y_train <- y_train[!(index(y_train) %in% influentials)]
removed_boxcox_y_train <- boxcox_y_train[!(index(boxcox_y_train) %in% influentials)]


removed_lassoAIC_linear_model <- lm(formula = removed_boxcox_y_train ~ ., data = removed_lassoAIC_X_train)

results <- summary(removed_lassoAIC_linear_model)
cat("R^2:", results$r.squared, "\n")
cat("Max beta's std. error:", max(results$coefficients[,2]), "\n")
cat("Mean beta's std. error:", mean(results$coefficients[,2]), "\n")

check_assumptions(removed_lassoAIC_linear_model, fitted(removed_lassoAIC_linear_model))
````
```



Normal Q-Q Plot



$\Rightarrow$



# Improvements

- Can do a better job in encoding
  - Huge datasets mean common assumptions tests like Shapiro-Wilk don't work in R
- => Assumptions validation is subjective
- Never forget to add training and testing result again! (4100~ training error, testing error ~ 4900)

***Thank you for listening!***